

1. One scene, one word apart




put the bowl on the PLATE


put the bowl on the STOVE

identical pixels; the action difference
is attributable to that word

2. Vision does not override



3. Untrained pairings degrade



still tracks the command, so not a fixed lookup